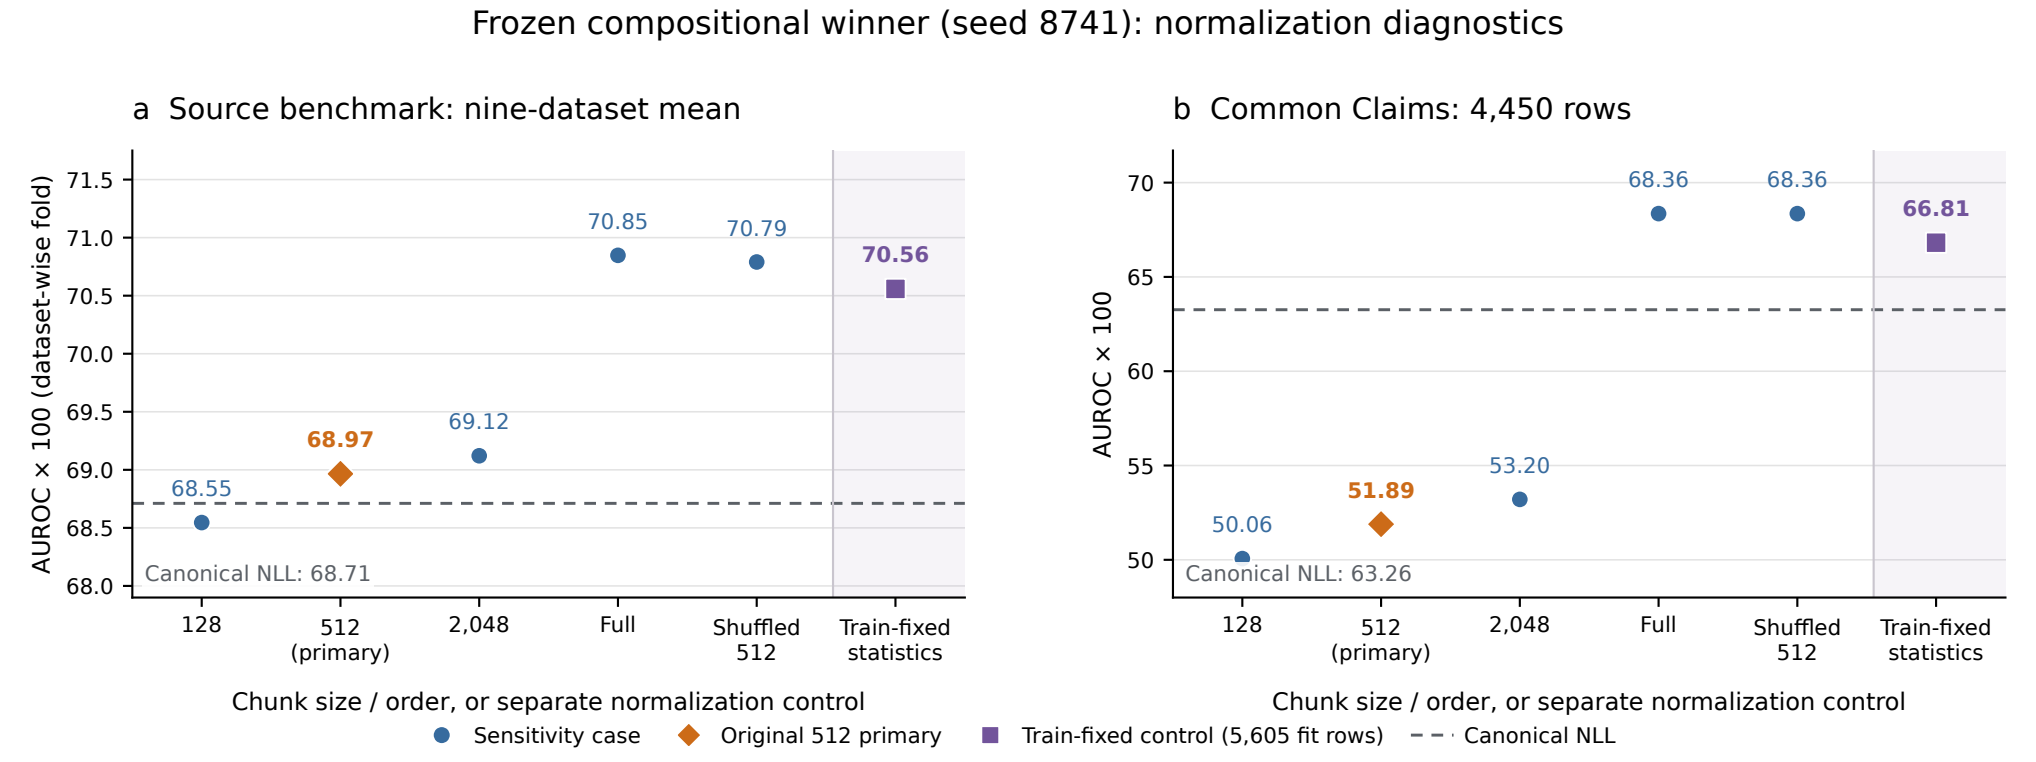


Common Claims order: 2,225 label-0 rows followed by 2,225 label-1 rows. No evaluation labels enter normalization or shuffling.

Original training, validation and benchmark scoring all used 512-row chunks. Other settings are diagnostics; the original primary is retained.

Point estimates only. The train-fixed result is a separate exploratory control designed after test viewing; all fixed cases are shown.